

# CSI62

## Operating Systems and Systems Programming

### Lecture 2

## Introduction to the Process


August 29<sup>th</sup>, 2016

Prof. Anthony D. Joseph

<http://cs162.eecs.Berkeley.edu>

## Recall: What is an operating system?

- Special layer of software that provides application software access to hardware resources
  - Convenient abstraction of complex hardware devices
  - Protected access to shared resources
  - Security and authentication
  - Communication amongst logical entities



8/29/16

Joseph CSI62 ©UCB Fall 2016

Lec 2.2

## Review: What is an Operating System?



- Referee
  - Manage sharing of resources, Protection, Isolation
    - » Resource allocation, isolation, communication



- Illusionist
  - Provide clean, easy to use abstractions of physical resources
    - » Infinite memory, dedicated machine
    - » Higher level objects: files, users, messages
    - » Masking limitations, virtualization



- Glue
  - Common services
    - » Storage, Window system, Networking
    - » Sharing, Authorization
    - » Look and feel


8/29/16

Joseph CSI62 ©UCB Fall 2016

Lec 2.3

## Review: Increasing Software Complexity

Millions of lines of source code



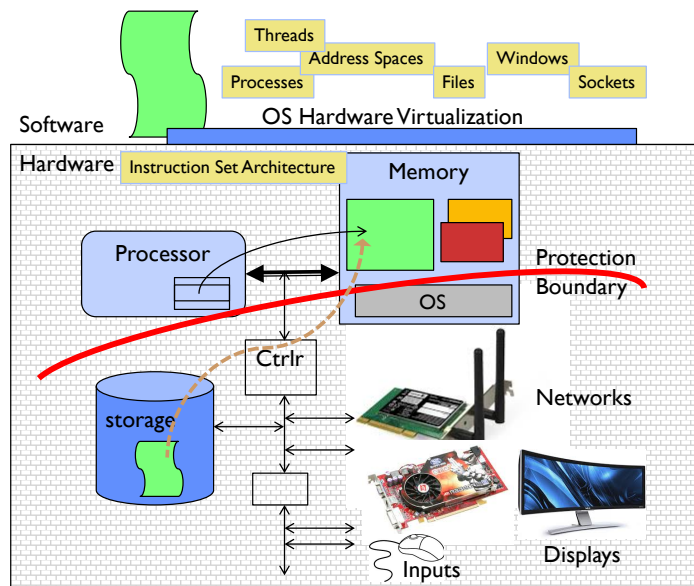
From MIT's 6.033 course

8/29/16

Joseph CSI62 ©UCB Fall 2016

Lec 2.4

## Recall: Loading



8/29/16

Joseph CSI62 ©UCB Fall 2016

Lec 2.5

## Very Brief History of OS

- Several Distinct Phases:
  - Hardware Expensive, Humans Cheap
    - » Eniac, ... Multics
  - Hardware Cheaper, Humans Expensive
    - » PCs, Workstations, Rise of GUIs
  - Hardware Really Cheap, Humans Really Expensive
    - » Ubiquitous devices, Widespread networking



"I think there is a world market for maybe five computers." – Thomas Watson, chairman of IBM, 1943

8/29/16

Joseph CSI62 ©UCB Fall 2016

Lec 2.6

## Very Brief History of OS

- Several Distinct Phases:
  - Hardware Expensive, Humans Cheap
    - » Eniac, ... Multics
  - Hardware Cheaper, Humans Expensive
    - » PCs, Workstations, Rise of GUIs
  - Hardware Really Cheap, Humans Really Expensive
    - » Ubiquitous devices, Widespread networking



Thomas Watson was often called "the worlds greatest salesman" by the time of his death in 1956



8/29/16

Joseph CSI62 ©UCB Fall 2016

Lec 2.7

## Very Brief History of OS

- Several Distinct Phases:
  - Hardware Expensive, Humans Cheap
    - » Eniac, ... Multics
  - Hardware Cheaper, Humans Expensive
    - » PCs, Workstations, Rise of GUIs
  - Hardware Really Cheap, Humans Really Expensive
    - » Ubiquitous devices, Widespread networking



8/29/16

Joseph CSI62 ©UCB Fall 2016

Lec 2.8

## Very Brief History of OS

- Several Distinct Phases:
  - Hardware Expensive, Humans Cheap
    - » Eniac, ... Multics
  - Hardware Cheaper, Humans Expensive
    - » PCs, Workstations, Rise of GUIs
  - Hardware Really Cheap, Humans Really Expensive
    - » Ubiquitous devices, Widespread networking
- Rapid Change in Hardware Leads to changing OS
  - Batch ⇒ Multiprogramming ⇒ Timesharing ⇒ Graphical UI ⇒ Ubiquitous Devices
  - Gradual Migration of Features into Smaller Machines
- Situation today is much like the late 60s
  - Small OS: 100K lines/Large: 10M lines (5M browser!)
  - 100-1000 people-years

8/29/16

Joseph CSI 62 ©UCB Fall 2016

Lec 2.9

## OS Archaeology

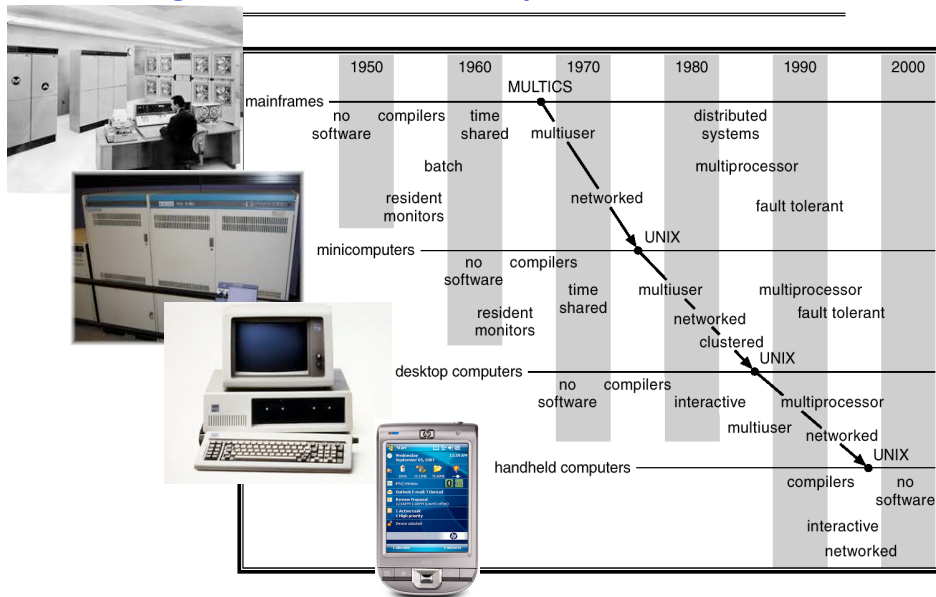
- Because of the cost of developing an OS from scratch, most modern OSes have a long lineage:
- Multics → AT&T Unix → BSD Unix → Ultrix, SunOS, NetBSD,...
- Mach (micro-kernel) + BSD → NextStep → XNU → Apple OS X, iPhone iOS
- MINIX → Linux → Android OS, Chrome OS, RedHat, Ubuntu, Fedora, Debian, Suse,...
- CP/M → QDOS → MS-DOS → Windows 3.1 → NT → 95 → 98 → 2000 → XP → Vista → 7 → 8 → 10 → phone → ...

8/29/16

Joseph CSI 62 ©UCB Fall 2016

Lec 2.10

## Migration of OS Concepts and Features



8/29/16

Joseph CSI 62 ©UCB Fall 2016

Lec 2.11

## Today: Four Fundamental OS Concepts

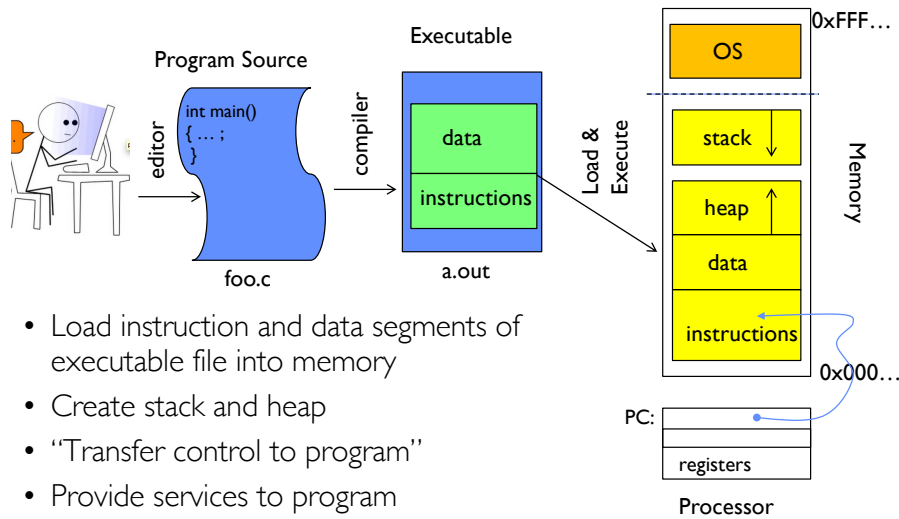
- Thread
  - Single unique execution context
  - Program Counter, Registers, Execution Flags, Stack
- Address Space with Translation
  - Programs execute in an *address space* that is distinct from the memory space of the physical machine
- Process
  - An instance of an executing program is a *process* consisting of an *address space* and one or more *threads of control*
- Dual Mode operation/Protection
  - Only the “system” has the ability to access certain resources
  - The OS and the hardware are protected from user programs and user programs are isolated from one another by *controlling the translation* from program virtual addresses to machine physical addresses

8/29/16

Joseph CSI 62 ©UCB Fall 2016

Lec 2.12

## OS Bottom Line: Run Programs



- Load instruction and data segments of executable file into memory
- Create stack and heap
- “Transfer control to program”
- Provide services to program
- While protecting OS and program

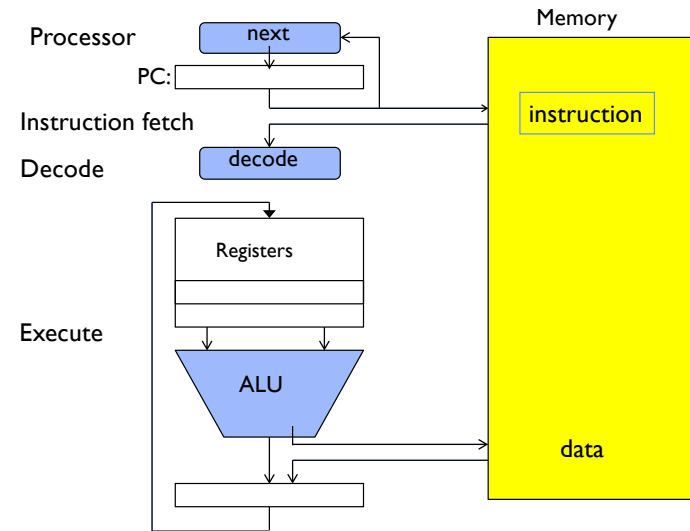
8/29/16

Joseph CSI62 ©UCB Fall 2016

Lec 2.13

## Recall (61B): Instruction Fetch/Decode/Execute

The instruction cycle

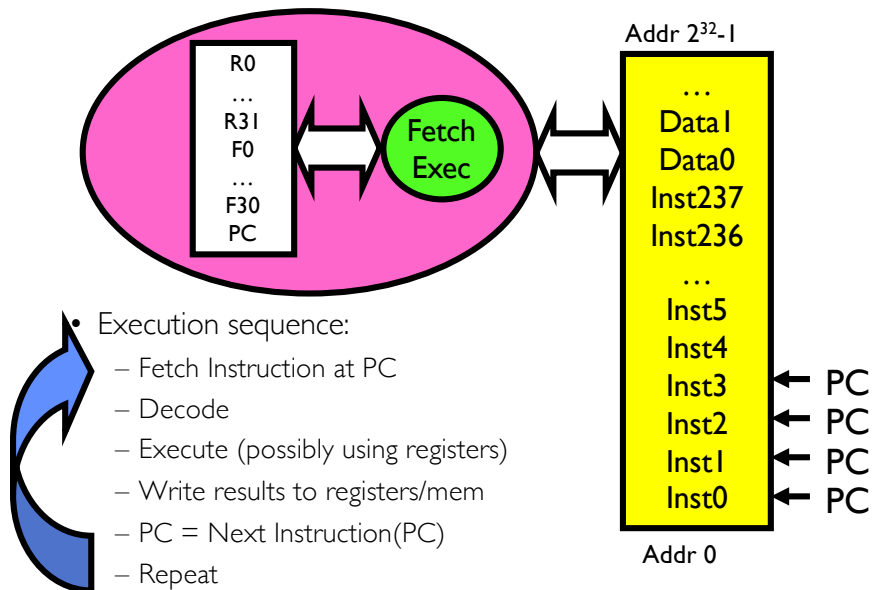


8/29/16

Joseph CSI62 ©UCB Fall 2016

Lec 2.14

## Recall (61C): What happens during program execution?



- Certain registers hold the *context* of thread
  - Stack pointer holds the address of the top of stack
    - » Other conventions: Frame Pointer, Heap Pointer, Data
  - May be defined by the instruction set architecture or by compiler conventions
- Thread: Single unique execution context
  - Program Counter, Registers, Execution Flags, Stack
- A thread is executing on a processor when it is resident in the processor registers.
- PC register holds the address of executing instruction in the thread.
- Registers hold the root state of the thread.
  - The rest is “in memory”

8/29/16

Joseph CSI62 ©UCB Fall 2016

Lec 2.15


8/29/16

Joseph CSI62 ©UCB Fall 2016

Lec 2.16

## Second OS Concept: Program's Address Space

- Address space  $\Rightarrow$  the set of accessible addresses + state associated with them:
  - For a 32-bit processor there are  $2^{32} = 4$  billion addresses
- What happens when you read or write to an address?
  - Perhaps Nothing
  - Perhaps acts like regular memory
  - Perhaps ignores writes
    - » (Memory-mapped I/O)
  - Perhaps causes exception (fault)

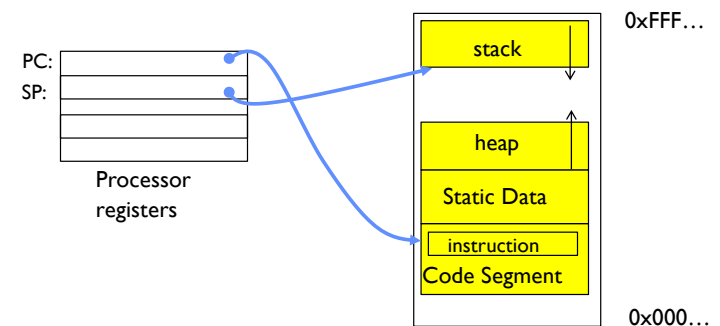


8/29/16

Joseph CSI 62 ©UCB Fall 2016

Lec 2.17

## Address Space: In a Picture



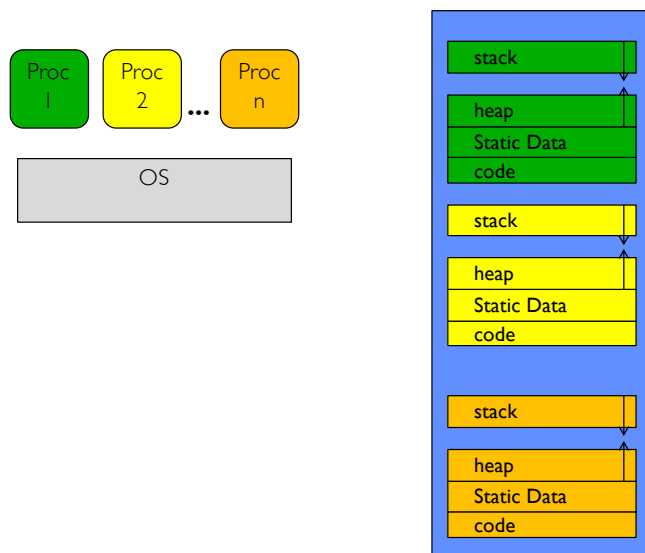
- What's in the code Segment? Static Data Segment?
- What's in the Stack Segment?
  - How is it allocated? How big is it?
- What's in the Heap Segment?
  - How is it allocated? How big?

8/29/16

Joseph CSI 62 ©UCB Fall 2016

Lec 2.18

## Multiprogramming - Multiple Threads of Control



8/29/16

Joseph CSI 62 ©UCB Fall 2016

Lec 2.19

## Administrivia: Getting started


- Start homework 0 immediately  $\Rightarrow$  **Due next Monday!**
  - cs162-xx account, Github account, registration survey
  - Vagrant and VirtualBox – VM environment for the course
    - » Consistent, managed environment on your machine
  - Get familiar with all the cs162 tools, submit to autograder via git
  - Homework slip days: **You have 3 slip days**
- Next Friday (9/9) is early drop day! Very hard to drop afterwards...**
- Should be going to section already!
- Group sign up form will be out after drop deadline
  - Work on finding groups ASAP: 4 people in a group!
  - Try to attend either same section or 2 sections by same TA

8/29/16

Joseph CSI 62 ©UCB Fall 2016

Lec 2.20

## How can we give the illusion of multiple processors?



- Assume a single processor. How do we provide the illusion of multiple processors?
  - Multiplex in time!
- Each virtual “CPU” needs a structure to hold:
  - Program Counter (PC), Stack Pointer (SP)
  - Registers (Integer, Floating point, others...?)
- How switch from one virtual CPU to the next?
  - Save PC, SP, and registers in current state block
  - Load PC, SP, and registers from new state block
- What triggers switch?
  - Timer, voluntary yield, I/O, other things

8/29/16

Joseph CSI62 ©UCB Fall 2016

Lec 2.21

## The Basic Problem of Concurrency

- The basic problem of concurrency involves resources:
  - Hardware: single CPU, single DRAM, single I/O devices
  - Multiprogramming API: processes think they have exclusive access to shared resources
- OS has to coordinate all activity
  - Multiple processes, I/O interrupts, ...
  - How can it keep all these things straight?
- Basic Idea: Use Virtual Machine abstraction
  - Simple machine abstraction for processes
  - Multiplex these abstract machines
- Dijkstra did this for the “THE system”
  - Few thousand lines vs 1 million lines in OS 360 (1K bugs)

8/29/16

Joseph CSI62 ©UCB Fall 2016

Lec 2.22

## Properties of this simple multiprogramming technique

- All virtual CPUs share same non-CPU resources
  - I/O devices the same
  - Memory the same
- Consequence of sharing:
  - Each thread can access the data of every other thread (good for sharing, bad for protection)
  - Threads can share instructions (good for sharing, bad for protection)
  - Can threads overwrite OS functions?
- This (unprotected) model is common in:
  - Embedded applications
  - Windows 3.1/Early Macintosh (switch only with yield)
  - Windows 95—ME (switch with both yield and timer)

8/29/16

Joseph CSI62 ©UCB Fall 2016

Lec 2.23

## Protection

- Operating System must protect itself from user programs
  - Reliability: compromising the operating system generally causes it to crash
  - Security: limit the scope of what processes can do
  - Privacy: limit each process to the data it is permitted to access
  - Fairness: each should be limited to its appropriate share of system resources (CPU time, memory, I/O, etc)
- It must protect User programs from one another
- Primary Mechanism: limit the translation from program address space to physical memory space
  - Can only touch what is mapped into process *address space*
- Additional Mechanisms:
  - Privileged instructions, in/out instructions, special registers
  - syscall processing, subsystem implementation
  - » (e.g., file access rights, etc)

8/29/16

Joseph CSI62 ©UCB Fall 2016

Lec 2.24

## Third OS Concept: Process

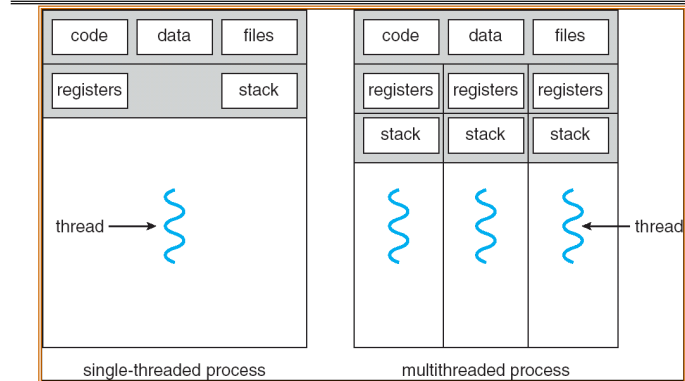
- **Process:** execution environment with Restricted Rights
  - Address Space with One or More Threads
  - Owns memory (address space)
  - Owns file descriptors, file system context, ...
  - Encapsulate one or more threads sharing process resources
- Why **processes**?
  - Protected from each other!
  - OS Protected from them
  - Processes provides memory protection
  - Threads more efficient than processes (later)
- Fundamental tradeoff between protection and efficiency
  - Communication easier *within* a process
  - Communication harder *between* processes
- Application instance consists of one or more processes

8/29/16

Joseph CSI 62 ©UCB Fall 2016

Lec 2.25

## Single and Multithreaded Processes



- Threads encapsulate concurrency: “Active” component
- Address spaces encapsulate protection: “Passive” part
  - Keeps buggy program from trashing the system
- Why have multiple threads per address space?

8/29/16

Joseph CSI 62 ©UCB Fall 2016

Lec 2.26

## Fourth OS Concept: Dual Mode Operation

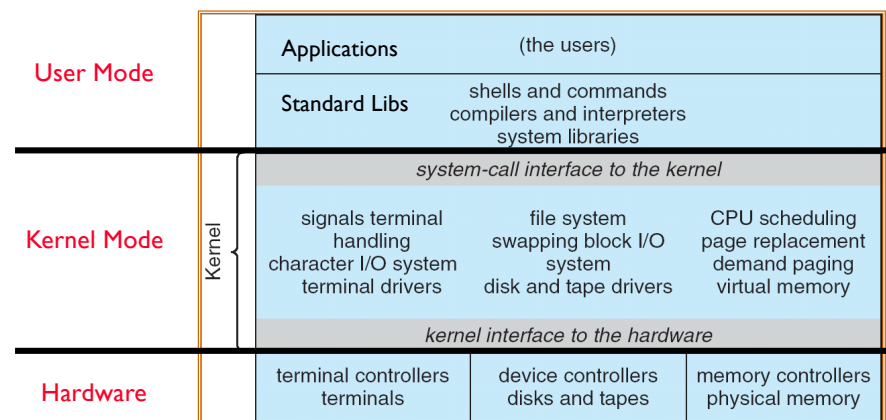
- **Hardware** provides at least two modes:
  - “Kernel” mode (or “supervisor” or “protected”)
  - “User” mode: Normal programs executed
- What is needed in the hardware to support “dual mode” operation?
  - a bit of state (user/system mode bit)
  - Certain operations / actions only permitted in system/kernel mode
    - » In user mode they fail or trap
  - User→Kernel transition sets system mode AND saves the user PC
    - » Operating system code carefully puts aside user state then performs the necessary operations
  - Kernel→User transition clears system mode AND restores appropriate user PC
    - » return-from-interrupt

8/29/16

Joseph CSI 62 ©UCB Fall 2016

Lec 2.27

## For example: UNIX System Structure




8/29/16

Joseph CSI 62 ©UCB Fall 2016

Lec 2.28

## User/Kernel (Privileged) Mode

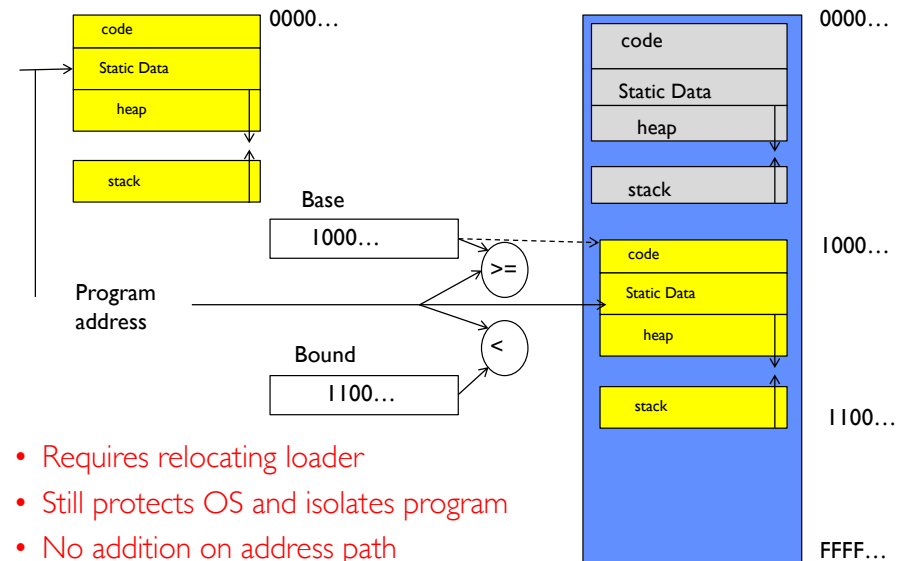


8/29/16

Joseph CSI 62 ©UCB Fall 2016

Lec 2.29

## Simple Protection: Base and Bound (B&B)



- Requires relocating loader
- Still protects OS and isolates program
- No addition on address path


8/29/16

Joseph CSI 62 ©UCB Fall 2016

Lec 2.30

## Another idea: Address Space Translation

- Program operates in an address space that is distinct from the physical memory space of the machine

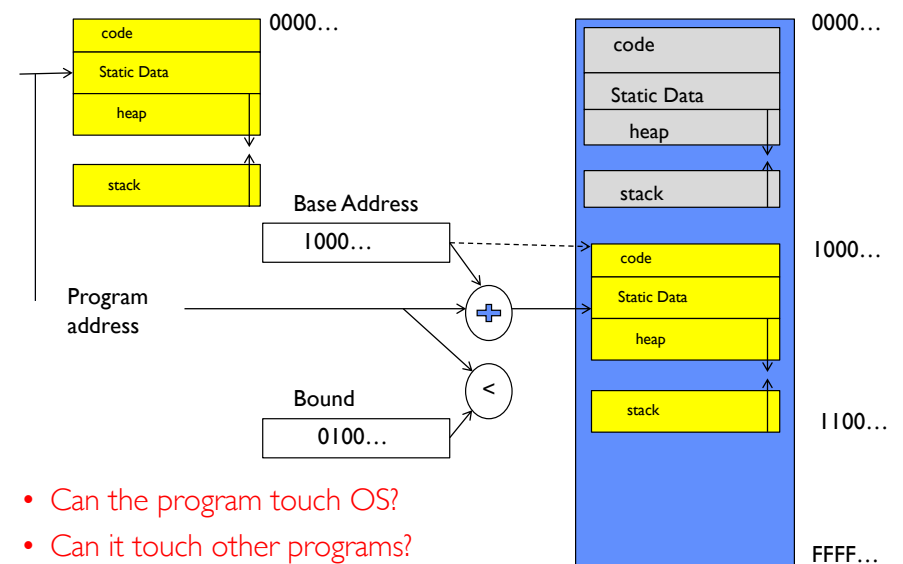


8/29/16

Joseph CSI 62 ©UCB Fall 2016

Lec 2.31

## A simple address translation with Base and Bound



- Can the program touch OS?
- Can it touch other programs?

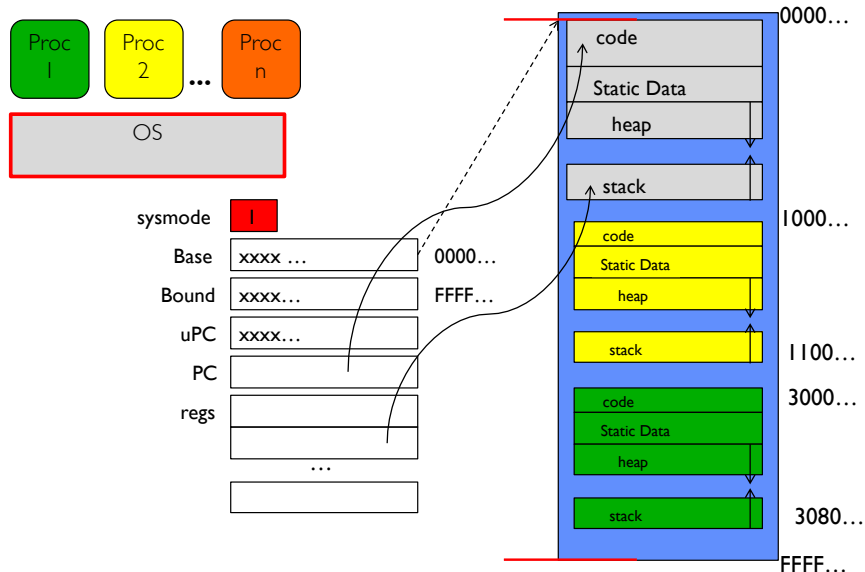
8/29/16

Joseph CSI 62 ©UCB Fall 2016

Lec 2.32



## Tying it together: Simple B&B: OS loads process

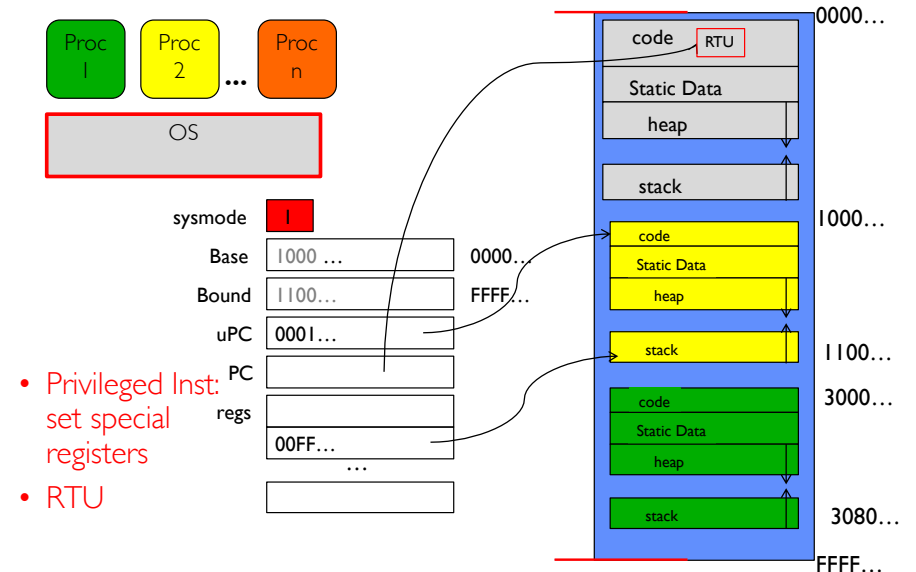


8/29/16

Joseph CSI62 ©UCB Fall 2016

Lec 2.33

## Simple B&B: OS gets ready to execute process



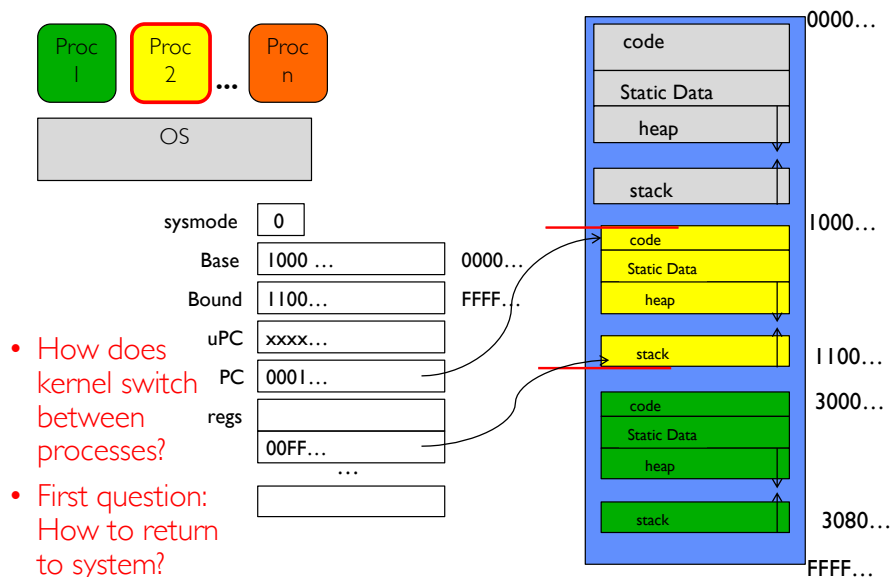
- Privileged Inst: set special registers
- RTU

8/29/16

Joseph CSI62 ©UCB Fall 2016

Lec 2.34

## Simple B&B: User Code Running



- How does kernel switch between processes?
- First question: How to return to system?

8/29/16

Joseph CSI62 ©UCB Fall 2016

Lec 2.35

## 3 types of Mode Transfer

- Syscall
  - Process requests a system service, e.g., exit
  - Like a function call, but "outside" the process
  - Does not have the address of the system function to call
  - Like a Remote Procedure Call (RPC) – for later
  - Marshall the syscall id and args in registers and exec syscall
- Interrupt
  - External asynchronous event triggers context switch
  - e.g. Timer, I/O device
  - Independent of user process
- Trap or Exception
  - Internal synchronous event in process triggers context switch
  - e.g., Protection violation (segmentation fault), Divide by zero, ...
- All 3 are an UNPROGRAMMED CONTROL TRANSFER
  - Where does it go?

8/29/16

Joseph CSI62 ©UCB Fall 2016

Lec 2.36

## Administrivia (Cont'd)

- Joseph Office Hours: Mondays/Tuesdays 10-11 in 465 Soda
  - No office hours tomorrow 8/30
- Avoid private Piazza posts – others have same question
- Three Free Online Textbooks:
  - Click on “Resources” link for a list of “Online Textbooks”
  - Can read O'Reilly books for free as long as on campus or VPN
    - » One book on Git, two books on C
- Webcast: <https://CalCentral.Berkeley.edu/> (CalNet sign in)
  - Webcast is \*NOT\* a replacement for coming to class!

8/29/16

Joseph CSI62 ©UCB Fall 2016

Lec 2.37

## CS 162 Collaboration Policy

- ✓ Explaining a concept to someone in another group
- ✓ Discussing algorithms/testing strategies with other groups
- ✓ Helping debug someone else's code (in another group)
- ✓ Searching online for generic algorithms (e.g., hash table)
- ✗ Sharing code or test cases with another group
- ✗ Copying OR reading another group's code or test cases
- ✗ Copying OR reading online code or test cases from prior years

We compare all project submissions against prior year submissions and online solutions and will take actions (described on the course overview page) against offenders


8/29/16

Joseph CSI62 ©UCB Fall 2016

Lec 2.38

How do we get the system target address of the “unprogrammed control transfer?”

## Interrupt Vector



- Where else do you see this dispatch pattern?

8/29/16

Joseph CSI62 ©UCB Fall 2016

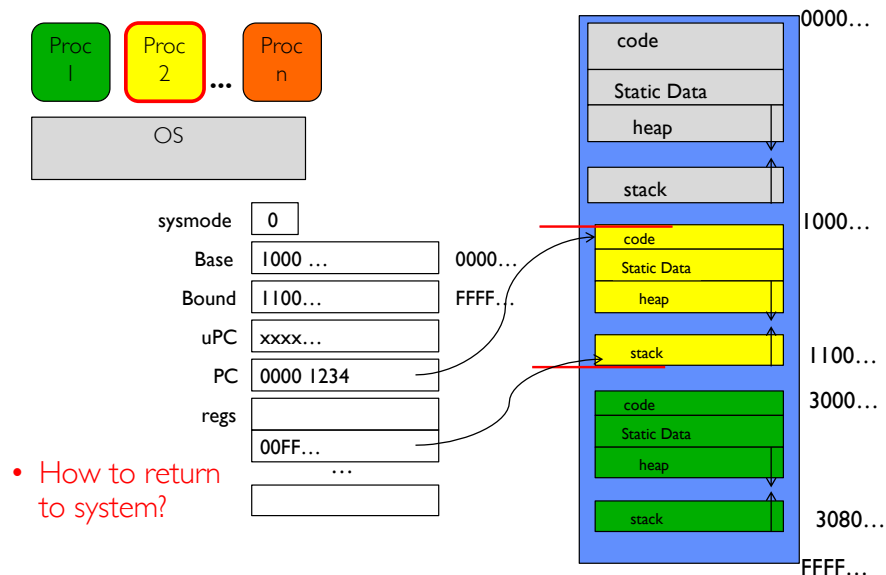
Lec 2.39

8/29/16

Joseph CSI62 ©UCB Fall 2016

Lec 2.40

## Simple B&B: User => Kernel



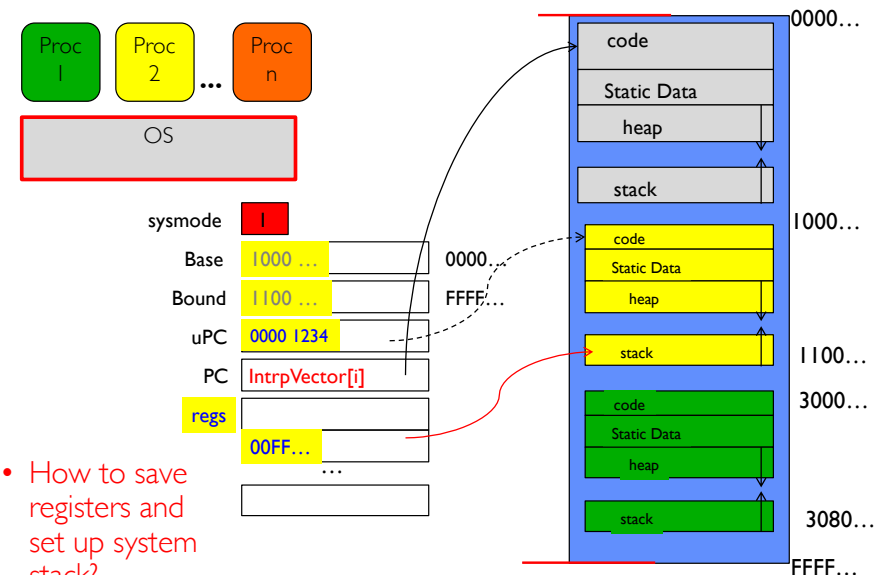
- How to return to system?

8/29/16

Joseph CSI62 ©UCB Fall 2016

Lec 2.41

## Simple B&B: Interrupt



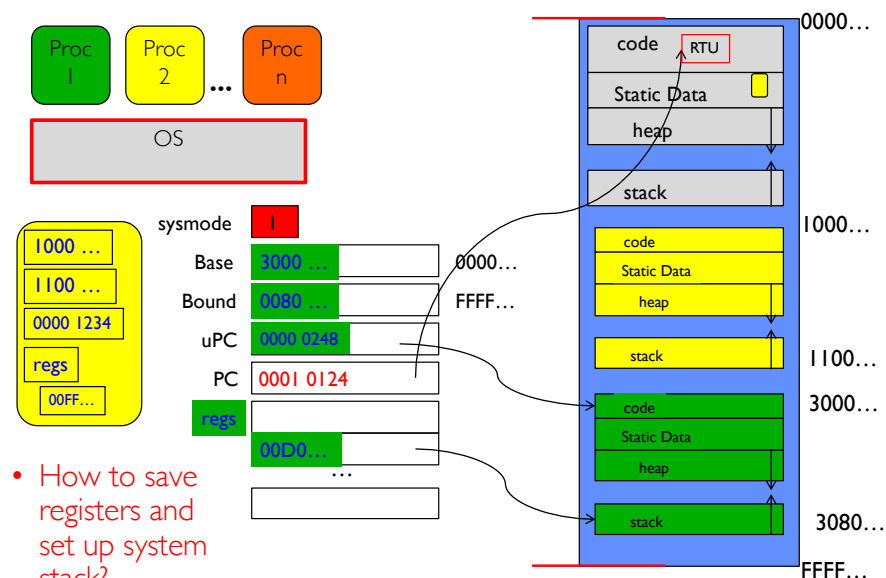
- How to save registers and set up system stack?

8/29/16

Joseph CSI62 ©UCB Fall 2016

Lec 2.42

## Simple B&B: Switch User Process



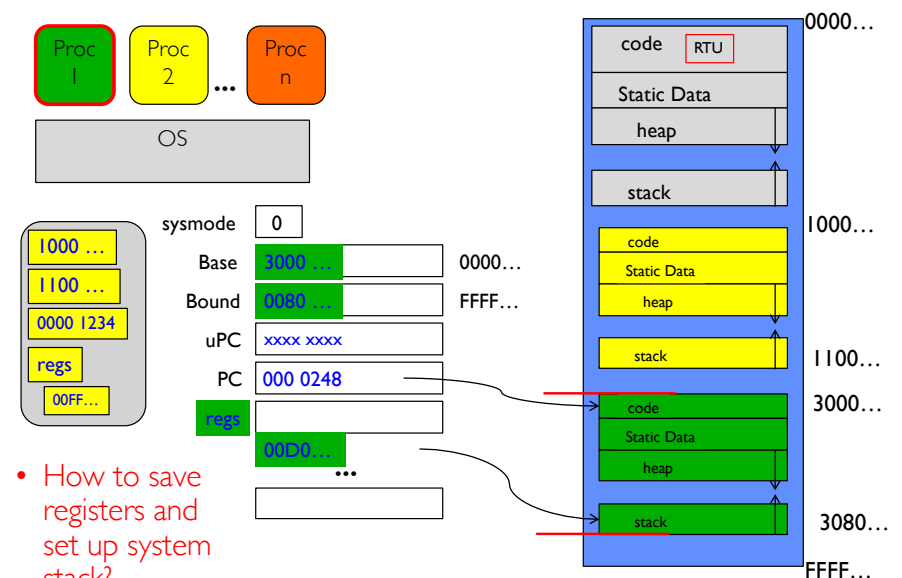
- How to save registers and set up system stack?

8/29/16

Joseph CSI62 ©UCB Fall 2016

Lec 2.43

## Simple B&B: "resume"



- How to save registers and set up system stack?

8/29/16

Joseph CSI62 ©UCB Fall 2016

Lec 2.44

## What's wrong with this simplistic address translation mechanism?

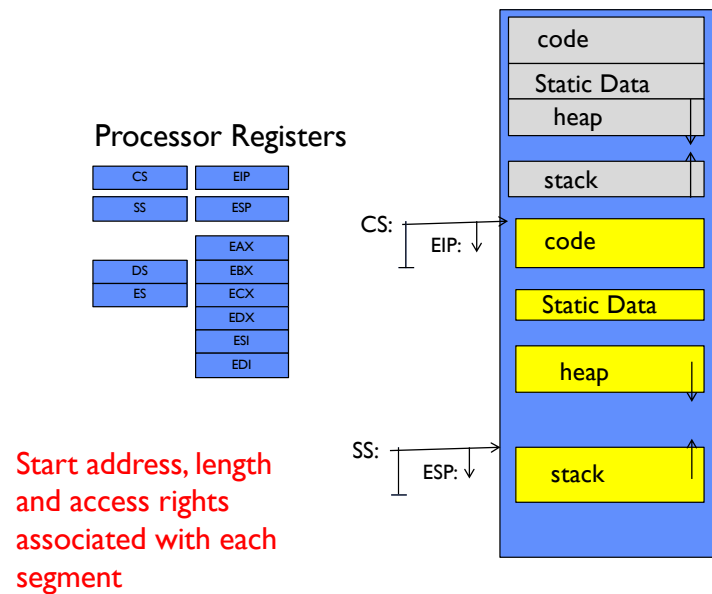
- Fragmentation:
  - Kernel has to somehow fit whole processes into contiguous block of memory
  - After a while, memory becomes fragmented!
- Sharing:
  - Very hard to share any data between Processes or between Process and Kernel
  - Simple segmentation

8/29/16

Joseph CSI 62 ©UCB Fall 2016

Lec 2.45

## x86 – segments and stacks



8/29/16

Joseph CSI 62 ©UCB Fall 2016

Lec 2.46

## Virtual Address Translation

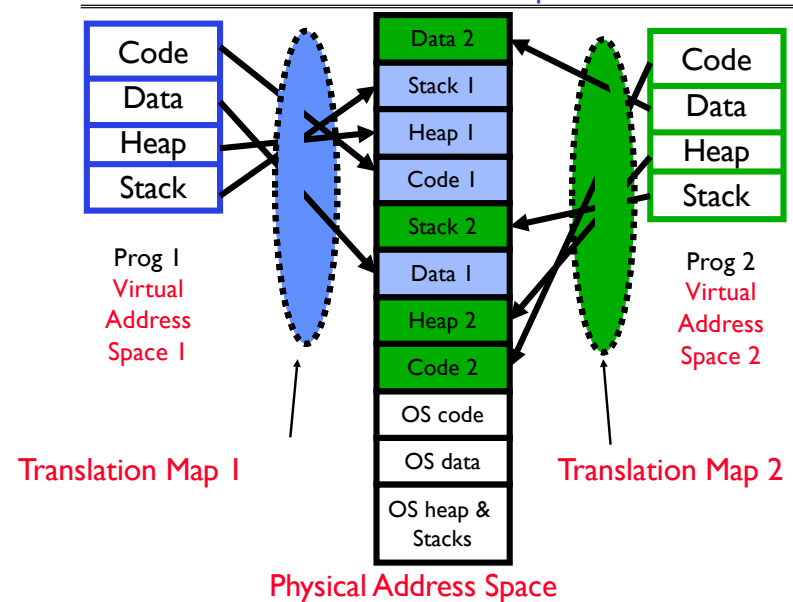
- Simpler, more useful schemes too!
- Give every process the illusion of its own **BIG FLAT ADDRESS SPACE**
  - Break it into pages
  - More on this later

8/29/16

Joseph CSI 62 ©UCB Fall 2016

Lec 2.47

## Providing Illusion of Separate Address Space: Load new Translation Map on Switch



8/29/16

Joseph CSI 62 ©UCB Fall 2016

Lec 2.48

## Running Many Programs ???

- We have the basic mechanism to
  - switch between user processes and the kernel,
  - the kernel can switch among user processes,
  - Protect OS from user processes and processes from each other
- Questions ???
- How do we decide which user process to run?
- How do we represent user processes in the OS?
- How do we pack up the process and set it aside?
- How do we get a stack and heap for the kernel?
- Aren't we wasting a lot of memory?
- ...

8/29/16

Joseph CSI 62 ©UCB Fall 2016

Lec 2.49

## Process Control Block

- Kernel represents each process as a process control block (PCB)
  - Status (running, ready, blocked, ...)
  - Register state (when not ready)
  - Process ID (PID), User, Executable, Priority, ...
  - Execution time, ...
  - Memory space, translation, ...
- Kernel Scheduler maintains a data structure containing the PCBs
- Scheduling algorithm selects the next one to run

8/29/16

Joseph CSI 62 ©UCB Fall 2016

Lec 2.50

## Scheduler

```

if ( readyProcesses(PCBs) ) {
    nextPCB = selectProcess(PCBs);
    run( nextPCB );
} else {
    run_idle_process();
}
  
```

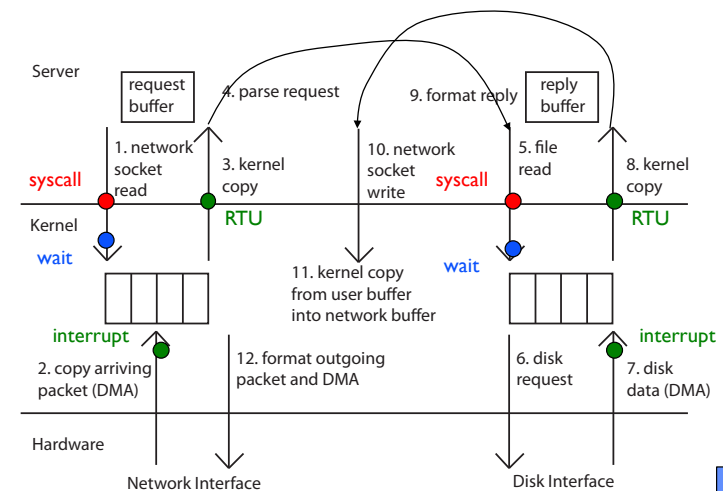
- Scheduling: Mechanism for deciding which processes/threads receive the CPU
- Lots of different scheduling policies provide ...
  - Fairness or
  - Realtime guarantees or
  - Latency optimization or ..

8/29/16

Joseph CSI 62 ©UCB Fall 2016

Lec 2.51

## Putting it together: web server



8/29/16

Joseph CSI 62 ©UCB Fall 2016

Lec 2.52

## Conclusion: Four fundamental OS concepts

- Thread
  - Single unique execution context
  - Program Counter, Registers, Execution Flags, Stack
- Address Space with Translation
  - Programs execute in an *address space* that is distinct from the memory space of the physical machine
- Process
  - An instance of an executing program is *a process consisting of an address space and one or more threads of control*
- Dual Mode operation/Protection
  - Only the “system” has the ability to access certain resources
  - The OS and the hardware are protected from user programs and user programs are isolated from one another by *controlling the translation* from program virtual addresses to machine physical addresses