

# CSI62

## Operating Systems and Systems Programming

### Lecture 3

## Processes (con't), Fork, Introduction to I/O

September 1<sup>st</sup>, 2016  
Prof. Anthony D. Joseph  
<http://cs162.eecs.Berkeley.edu>

## Recall: Four fundamental OS concepts

- Thread
  - Single unique execution context
  - Program Counter, Registers, Execution Flags, Stack
- Address Space w/ Translation
  - Programs execute in an *address space* that is distinct from the memory space of the physical machine
- Process
  - An instance of an executing program is a *process* consisting of an *address space* and one or more *threads of control*
- Dual Mode operation/Protection
  - Only the “system” has the ability to access certain resources
  - The OS and the hardware are protected from user programs and user programs are isolated from one another by *controlling the translation* from program virtual addresses to machine physical addresses

9/1/16

Joseph CSI62 ©UCB Fall 2016

Lec 3.2

## Process Control Block


- Kernel represents each process as a process control block (PCB)
  - Status (running, ready, blocked, ...)
  - Register state (when not ready)
  - Process ID (PID), User, Executable, Priority, ...
  - Execution time, ...
  - Memory space, translation, ...
- Kernel Scheduler maintains a data structure containing the PCBs
- Scheduling algorithm selects the next one to run

9/1/16

Joseph CSI62 ©UCB Fall 2016

Lec 3.3

## Recall: give the illusion of multiple processors?



- Assume a single processor. How do we provide the *illusion* of multiple processors?
  - Multiplex in time!
  - Multiple “virtual CPUs”
- Each virtual “CPU” needs a structure to hold:
  - Program Counter (PC), Stack Pointer (SP)
  - Registers (Integer, Floating point, others...?)
- How switch from one virtual CPU to the next?
  - Save PC, SP, and registers in current state block
  - Load PC, SP, and registers from new state block
- What triggers switch?
  - Timer, voluntary yield, I/O, other things


9/1/16

Joseph CSI62 ©UCB Fall 2016

Lec 3.4

## Simultaneous MultiThreading/Hyperthreading

- Hardware technique
  - Superscalar processors can execute multiple instructions that are independent
  - Hyperthreading duplicates register state to make a second “thread,” allowing more instructions to run
- Can schedule each thread as if were separate CPU
  - But, sub-linear speedup!
- Original technique called “Simultaneous Multithreading”
  - <http://www.cs.washington.edu/research/smt/index.html>
  - SPARC, Pentium 4/Xeon (“Hyperthreading”), Power 5



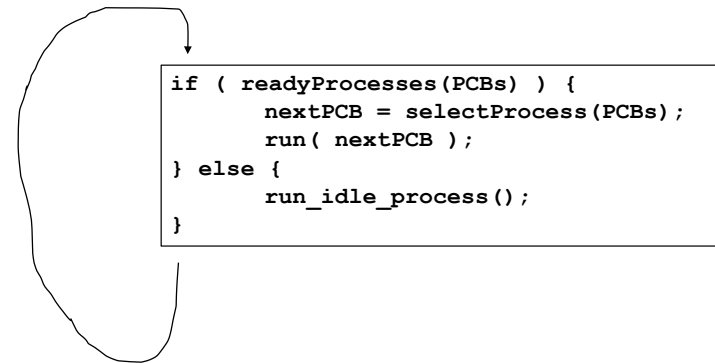
Colored blocks show instructions executed

9/1/16

Joseph CSI 62 ©UCB Fall 2016

Lec 3.5

## Scheduler



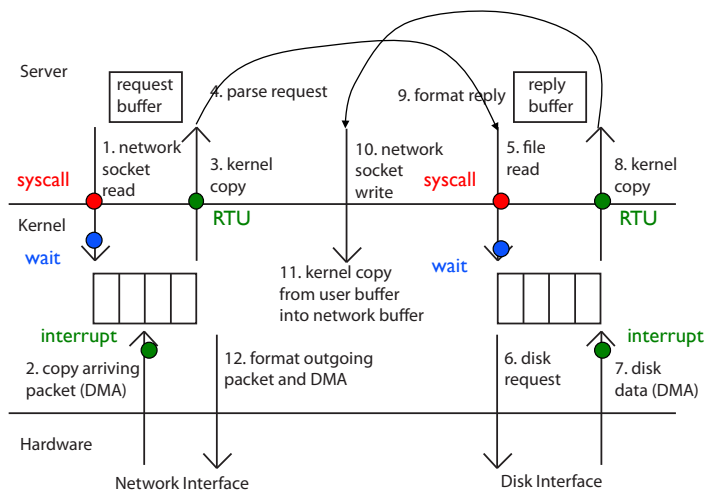
- Scheduling: Mechanism for deciding which processes/threads receive the CPU
- Lots of different scheduling policies provide ...
  - Fairness or
  - Realtime guarantees or
  - Latency optimization or ..

9/1/16

Joseph CSI 62 ©UCB Fall 2016

Lec 3.6

## Putting it together: web server



9/1/16

Joseph CSI 62 ©UCB Fall 2016

Lec 3.7

## Recall: 3 types of Kernel Mode Transfer


- Syscall
  - Process requests a system service, e.g., exit
  - Like a function call, but “outside” the process
  - Does not have the address of the system function to call
  - Like a Remote Procedure Call (RPC) – for later
  - Marshall the syscall ID and arguments in registers and execute syscall
- Interrupt
  - External asynchronous event triggers context switch
  - eg. Timer, I/O device
  - Independent of user process
- Trap or Exception
  - Internal synchronous event in process triggers context switch
  - e.g., Protection violation (segmentation fault), Divide by zero, ...

9/1/16

Joseph CSI 62 ©UCB Fall 2016

Lec 3.8

## Recall: User/Kernel (Privileged) Mode



9/1/16

Joseph CSI 62 ©UCB Fall 2016

Lec 3.9

## Implementing Safe Kernel Mode Transfers

- Important aspects:
  - Separate kernel stack
  - Controlled transfer into kernel (e.g. syscall table)
- Carefully constructed kernel code packs up the user process state and sets it aside
  - Details depend on the machine architecture
- Should be impossible for buggy or malicious user program to cause the kernel to corrupt itself

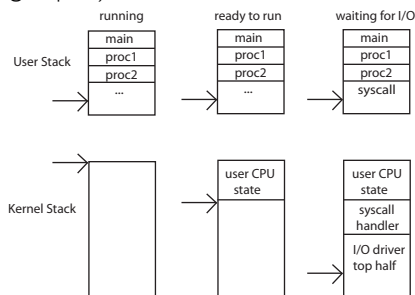
9/1/16

Joseph CSI 62 ©UCB Fall 2016

Lec 3.10

## Need for Separate Kernel Stacks

- Kernel needs space to work
- Cannot put anything on the user stack (Why?)
- Two-stack model
  - OS thread has interrupt stack (located in kernel memory) plus User stack (located in user memory)
  - Syscall handler copies user args to kernel space before invoking specific function (e.g., open)
  - Interrupts (???)

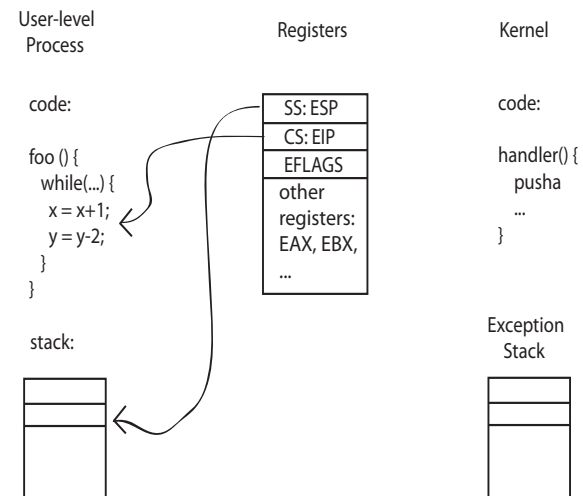


9/1/16

Joseph CSI 62 ©UCB Fall 2016

Lec 3.11

## Before

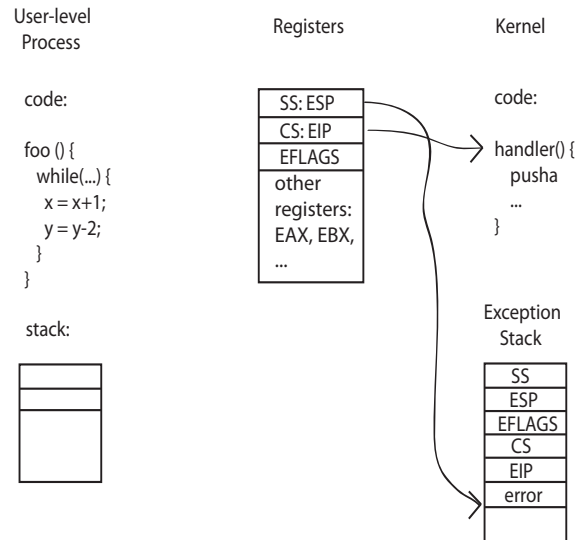


9/1/16

Joseph CSI 62 ©UCB Fall 2016

Lec 3.12

## During



9/1/16

Joseph CSI 62 ©UCB Fall 2016

Lec 3.13

## Kernel System Call Handler

- Vector through well-defined syscall entry points!
  - Table mapping system call number to handler
- Locate arguments
  - In registers or on user(!) stack
- Copy arguments
  - From user memory into kernel memory
  - Protect kernel from malicious code evading checks
- Validate arguments
  - Protect kernel from errors in user code
- Copy results back
  - into user memory

9/1/16

Joseph CSI 62 ©UCB Fall 2016

Lec 3.14

## Administrivia: Getting started

- We have a new discussion section!
  - 162-112 Th 6:30-7:30P 310 Soda
- Joseph Office Hours: Mondays/Tuesdays 11-12 in 465F Soda
- **THIS** Friday (9/2) is early drop day! Very hard to drop afterwards...
- Work on Homework 0 immediately  $\Rightarrow$  **Due on Monday!**
  - Get familiar with all the cs162 tools
  - Submit to autograder via git

9/1/16

Joseph CSI 62 ©UCB Fall 2016

Lec 3.15

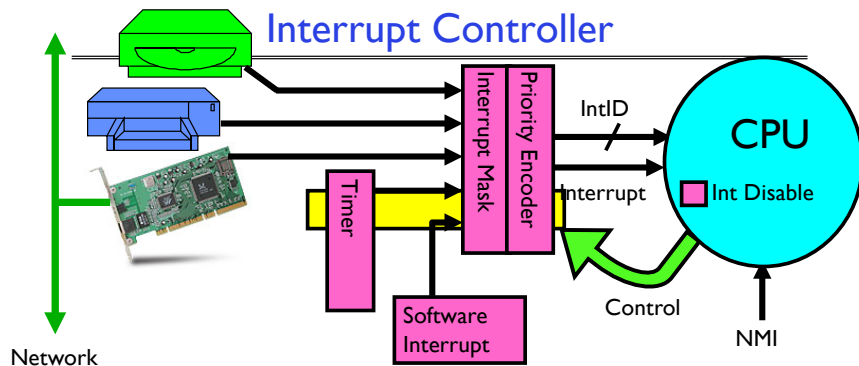
## Hardware support: Interrupt Control

- Interrupt processing not be visible to the user process:
  - Occurs between instructions, restarted transparently
  - No change to process state
  - What can be observed even with perfect interrupt processing?
- Interrupt Handler invoked with interrupts 'disabled'
  - Re-enabled upon completion
  - Non-blocking (run to completion, no waits)
  - Pack up in a queue and pass off to an OS thread for hard work
    - » wake up an existing OS thread
- OS kernel may enable/disable interrupts
  - On x86: CLI (disable interrupts), STI (enable)
  - Atomic section when select next process/thread to run
  - Atomic return from interrupt or syscall
- HW may have multiple levels of interrupt
  - Mask off (disable) certain interrupts, eg., lower priority
  - Certain Non-Maskable-Interrupts (NMI)
    - » e.g., kernel segmentation fault

9/1/16

Joseph CSI 62 ©UCB Fall 2016

Lec 3.17



- Interrupts invoked with interrupt lines from devices
- Interrupt controller chooses interrupt request to honor
  - Mask enables/disables interrupts
  - Priority encoder picks highest enabled interrupt
  - Software Interrupt Set/Cleared by Software
  - Interrupt identity specified with ID line
- CPU can disable all interrupts with internal flag
- Non-Maskable Interrupt line (NMI) can't be disabled

9/1/16

Joseph CSI 62 ©UCB Fall 2016

Lec 3.18

## How do we take interrupts safely?

- Interrupt vector
  - Limited number of entry points into kernel
- Kernel interrupt stack
  - Handler works regardless of state of user code
- Interrupt masking
  - Handler is non-blocking
- Atomic transfer of control
  - “Single instruction”-like to change:
    - » Program counter
    - » Stack pointer
    - » Memory protection
    - » Kernel/user mode
- Transparent restartable execution
  - User program does not know interrupt occurred

9/1/16

Joseph CSI 62 ©UCB Fall 2016

Lec 3.19

## Can a process create a process ?

- Yes! Unique identity of process is the “process ID” (or PID)
- Fork() system call creates a *copy* of current process with a new PID
- Return value from Fork(): integer
  - When > 0:
    - » Running in (original) **Parent** process
    - » return value is **pid** of new child
  - When = 0:
    - » Running in new **Child** process
  - When < 0:
    - » Error! Must handle somehow
    - » Running in original process
- All state of original process duplicated in both Parent and Child!
  - Memory, File Descriptors (next topic), etc...

9/1/16

Joseph CSI 62 ©UCB Fall 2016

Lec 3.20

## fork1.c

```
#include <stdlib.h>
#include <stdio.h>
#include <string.h>
#include <unistd.h>
#include <sys/types.h>

#define BUFSIZE 1024
int main(int argc, char *argv[])
{
    char buf[BUFSIZE];
    size_t readlen, writelen, slen;
    pid_t cpid, mypid;
    pid_t pid = getpid(); /* get current processes PID */
    printf("Parent pid: %d\n", pid);
    cpid = fork();
    if (cpid > 0) { /* Parent Process */
        mypid = getpid();
        printf("[%d] parent of [%d]\n", mypid, cpid);
    } else if (cpid == 0) { /* Child Process */
        mypid = getpid();
        printf("[%d] child\n", mypid);
    } else {
        perror("Fork failed");
        exit(1);
    }
    exit(0);
}
```

9/1/16

Joseph CSI 62 ©UCB Fall 2016

Lec 3.21

## UNIX Process Management

- UNIX **fork** – system call to create a copy of the current process, and start it running
  - No arguments!
- UNIX **exec** – system call to *change the program* being run by the current process
- UNIX **wait** – system call to wait for a process to finish
- UNIX **signal** – system call to send a notification to another process
- UNIX man pages: **fork(2)**, **exec(3)**, **wait(2)**, **signal(3)**

9/1/16

Joseph CSI 62 ©UCB Fall 2016

Lec 3.22

## fork2.c

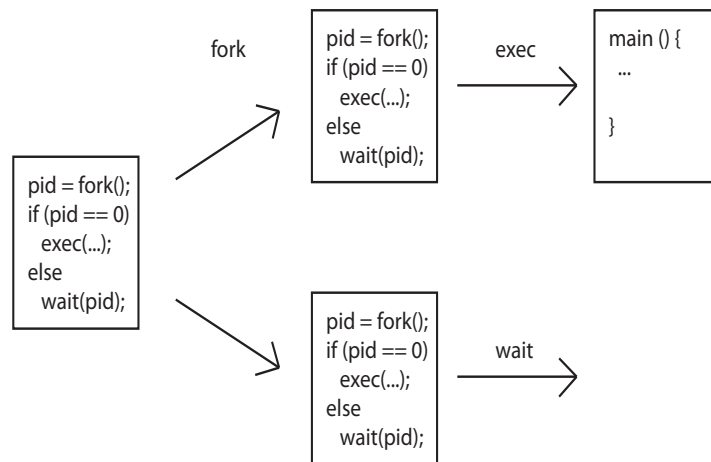
```
int status;
...
cpid = fork();
if (cpid > 0) {                /* Parent Process */
    mypid = getpid();
    printf("[%d] parent of [%d]\n", mypid, cpid);
    tcpid = wait(&status);
    printf("[%d] bye %d(%d)\n", mypid, tcpid, status);
} else if (cpid == 0) {        /* Child Process */
    mypid = getpid();
    printf("[%d] child\n", mypid);
}
...
```

9/1/16

Joseph CSI 62 ©UCB Fall 2016

Lec 3.23

## UNIX Process Management



9/1/16

Joseph CSI 62 ©UCB Fall 2016

Lec 3.24

## Administrivia (Con't)

- Participation: Attend section! Get to know your TA!
- Please use private Piazza posts *only* for student logistics issues
- Group sign up via autograder then TA form next week (after EDD)
  - Get finding groups of 4 people ASAP
  - Priority for same section; if cannot make this work, keep same TA

9/1/16

Joseph CSI 62 ©UCB Fall 2016

Lec 3.25

## BREAK

9/1/16

Joseph CSI 62 ©UCB Fall 2016

Lec 3.26

## Shell

- A shell is a job control system
  - Allows programmer to create and manage a set of programs to do some task
  - Windows, MacOS, Linux all have shells
- Example: to compile a C program

```
cc -c sourcefile1.c
cc -c sourcefile2.c
ln -o program sourcefile1.o sourcefile2.o
./program
```



9/1/16

Joseph CSI 62 ©UCB Fall 2016

Lec 3.27

## Signals – infloop.c

```
#include <stdlib.h>
#include <stdio.h>
#include <sys/types.h>

#include <unistd.h>
#include <signal.h>

void signal_callback_handler(int signum)
{
    printf("Caught signal %d - phew!\n", signum);
    exit(1);
}

int main() {
    signal(SIGINT, signal_callback_handler);

    while (1) {}
}
```

Got top?

9/1/16

Joseph CSI 62 ©UCB Fall 2016

Lec 3.28

## Process Races: fork3.c

```
int i;
cpid = fork();
if (cpid > 0) {
    mypid = getpid();
    printf("[%d] parent of [%d]\n", mypid, cpid);
    for (i=0; i<100; i++) {
        printf("[%d] parent: %d\n", mypid, i);
        // sleep(1);
    }
} else if (cpid == 0) {
    mypid = getpid();
    printf("[%d] child\n", mypid);
    for (i=0; i>-100; i--) {
        printf("[%d] child: %d\n", mypid, i);
        // sleep(1);
    }
}
```

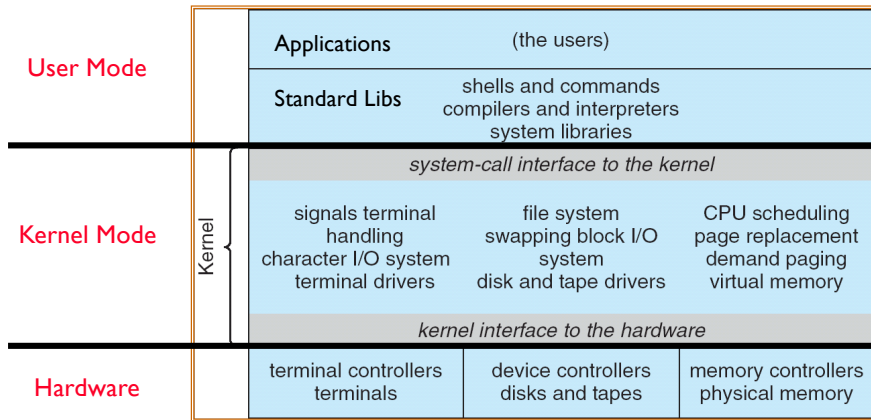
- Question: What does this program print?
- Does it change if you add in one of the sleep() statements?

9/1/16

Joseph CSI 62 ©UCB Fall 2016

Lec 3.29

## Recall: UNIX System Structure



9/1/16

Joseph CSI 62 ©UCB Fall 2016

Lec 3.30

## How Does the Kernel Provide Services?

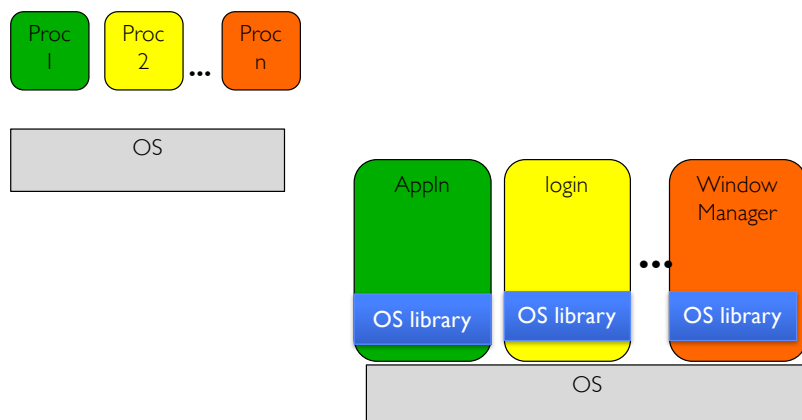
- You said that applications request services from the operating system via **syscall**, but ...
- I've been writing all sort of useful applications and I never ever saw a "syscall" !!!
- That's right.
- It was buried in the programming language runtime library (e.g., libc.a)
- ... Layering

9/1/16

Joseph CSI 62 ©UCB Fall 2016

Lec 3.31

## OS Run-Time Library

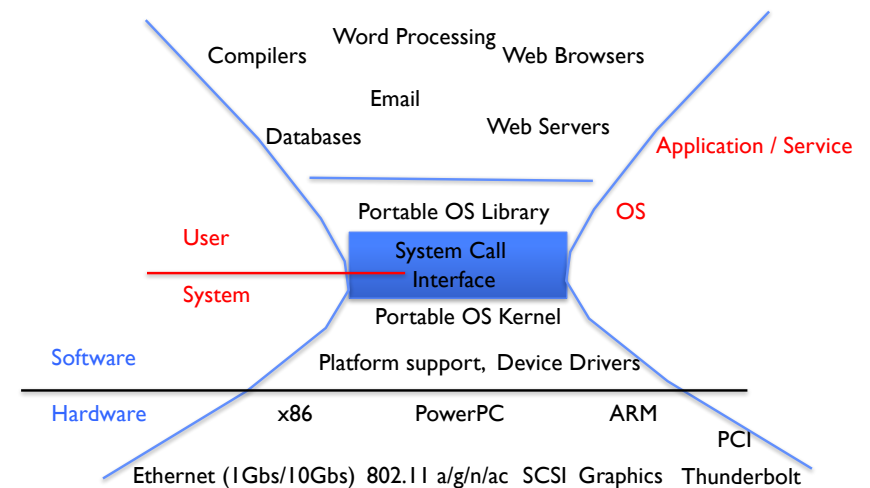


9/1/16

Joseph CSI 62 ©UCB Fall 2016

Lec 3.32

## A Kind of Narrow Waist



9/1/16

Joseph CSI 62 ©UCB Fall 2016

Lec 3.33



## Key Unix I/O Design Concepts

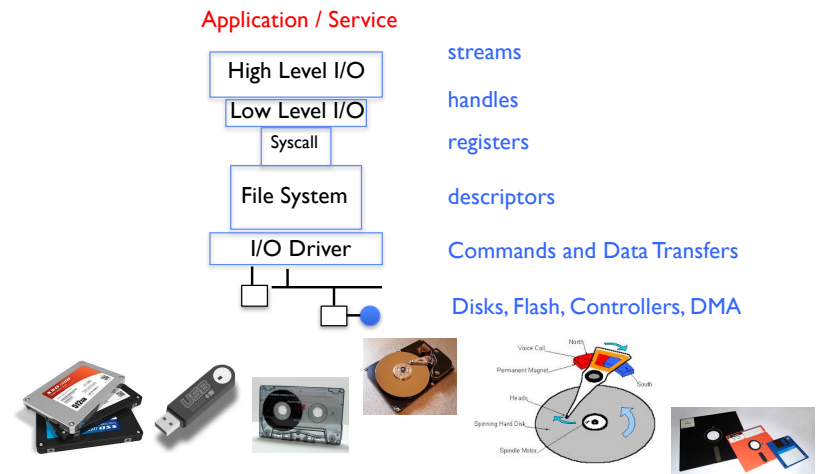
- Uniformity
  - file operations, device I/O, and interprocess communication through open, read/write, close
  - Allows simple composition of programs
    - » find | grep | wc ...
- Open before use
  - Provides opportunity for access control and arbitration
  - Sets up the underlying machinery, i.e., data structures
- Byte-oriented
  - Even if blocks are transferred, addressing is in bytes
- Kernel buffered reads
  - Streaming and block devices looks the same
  - read blocks process, yielding processor to other task
- Kernel buffered writes
  - Completion of out-going transfer decoupled from the application, allowing it to continue
- Explicit close

9/1/16

Joseph CSI 62 ©UCB Fall 2016

Lec 3.34

## I/O & Storage Layers



9/1/16

Joseph CSI 62 ©UCB Fall 2016

Lec 3.35

## The File System Abstraction

- High-level idea
  - Files live in hierarchical namespace of filenames
- File
  - Named collection of data in a file system
  - File data
    - » Text, binary, linearized objects
  - File Metadata: information about the file
    - » Size, Modification Time, Owner, Security info
    - » Basis for access control
- Directory
  - “Folder” containing files & Directories
  - Hierarchical (graphical) naming
    - » Path through the directory graph
    - » Uniquely identifies a file or directory
      - /home/ff/cs162/public\_html/fa16/index.html
  - Links and Volumes (later)

9/1/16

Joseph CSI 62 ©UCB Fall 2016

Lec 3.36

## Summary

- Process: execution environment with Restricted Rights
  - Address Space with One or More Threads
  - Owns memory (address space)
  - Owns file descriptors, file system context, ...
  - Encapsulate one or more threads sharing process resources
- Interrupts
  - Hardware mechanism for regaining control from user
  - Notification that events have occurred
  - User-level equivalent: Signals
- Native control of Process
  - Fork, Exec, Wait, Signal
- Basic Support for I/O
  - Standard interface: open, read, write, seek
  - Device drivers: customized interface to hardware

9/1/16

Joseph CSI 62 ©UCB Fall 2016

Lec 3.37